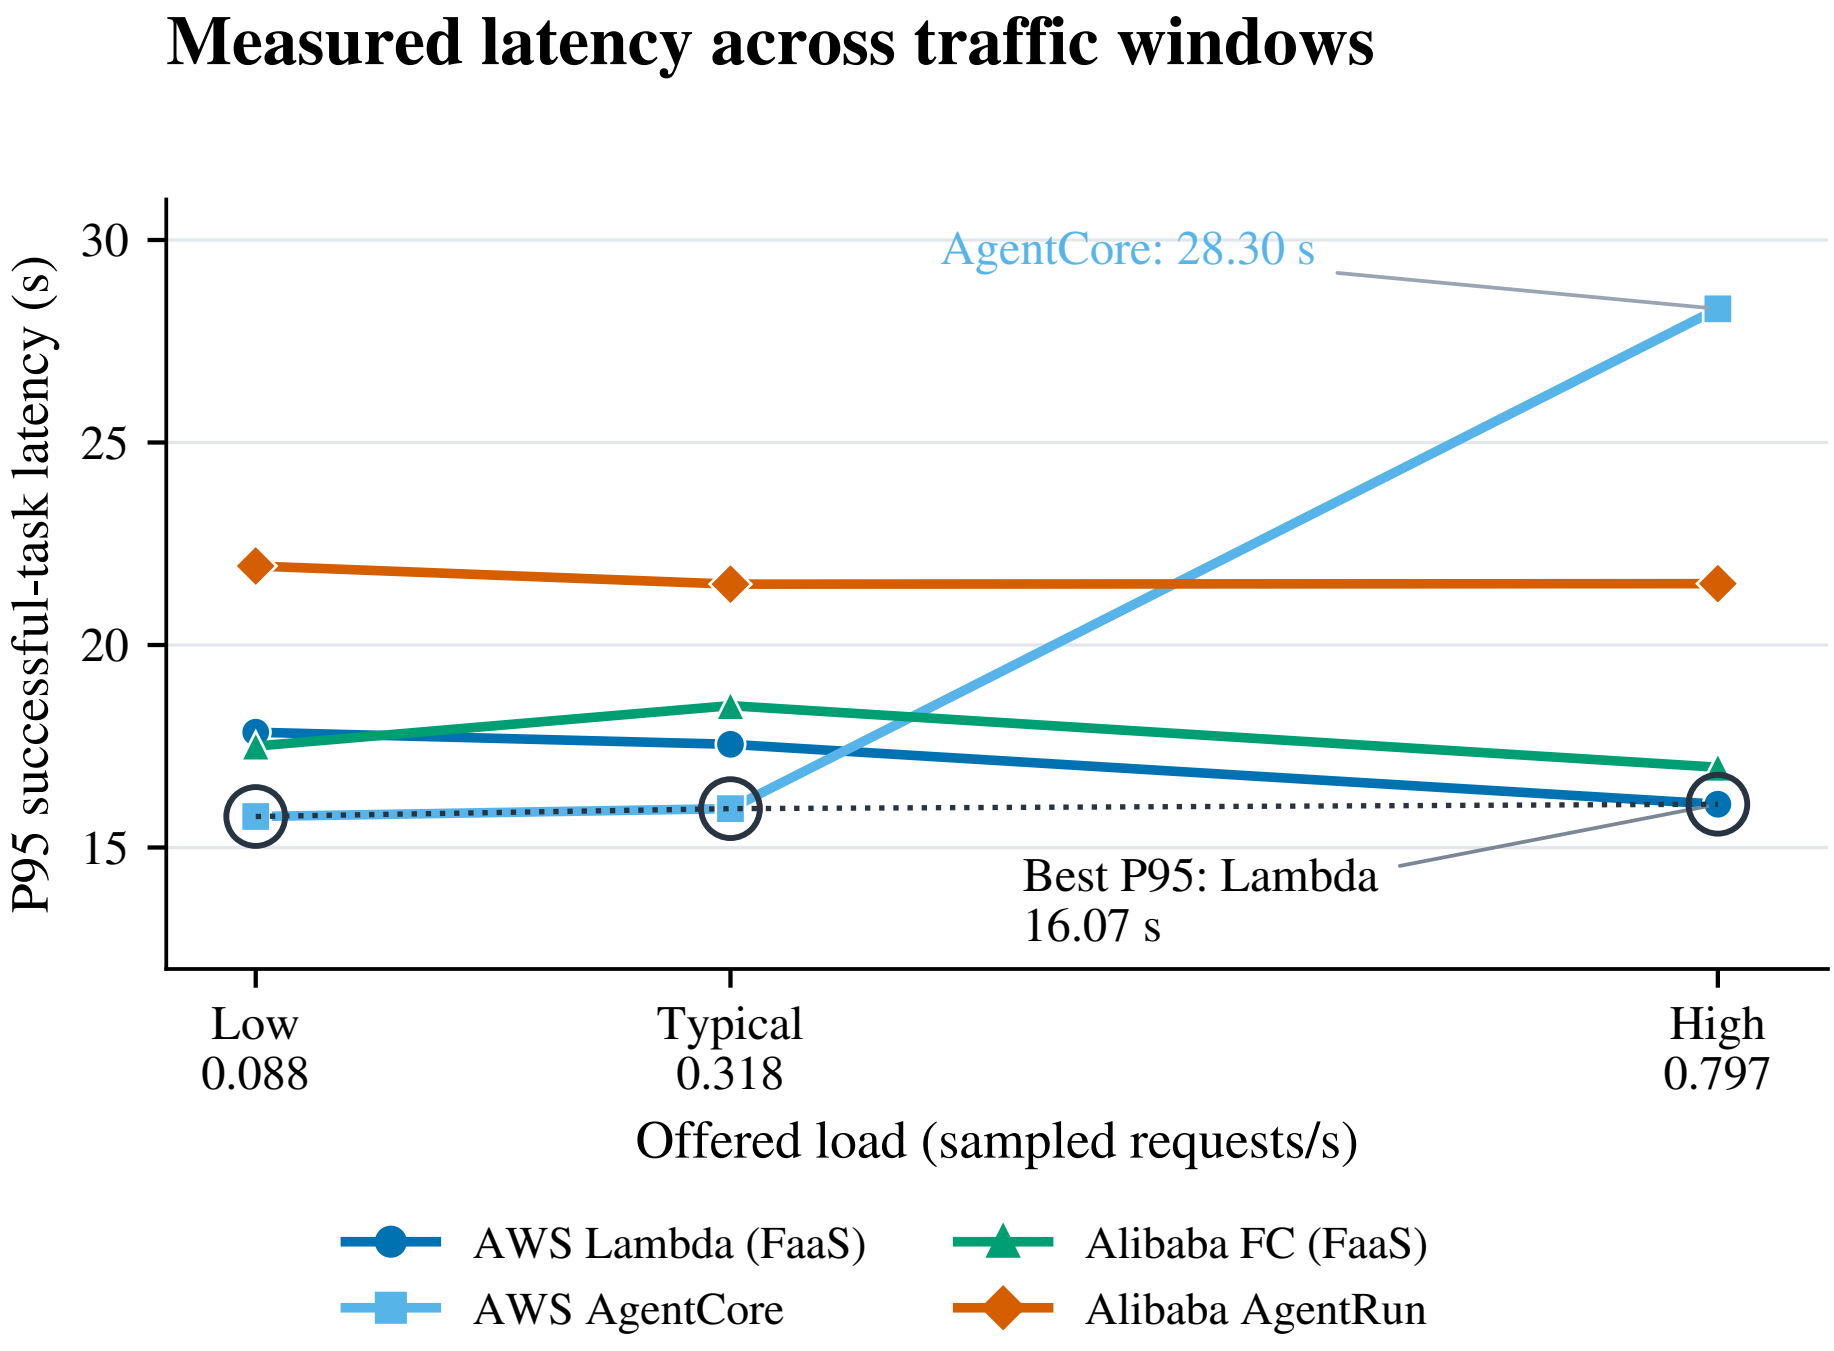
Measured latency across traffic windows
P95 successful-task latency (s)
30
25
20
15
AgentCore: 28.30 s
Best P95: Lambda
16.07 s
Low
0.088
Typical
0.318
High
0.797
Offered load (sampled requests/s)
AWS Lambda (FaaS)
Alibaba FC (FaaS)
AWS AgentCore
Alibaba AgentRun

Traffic-dependent deployment choices

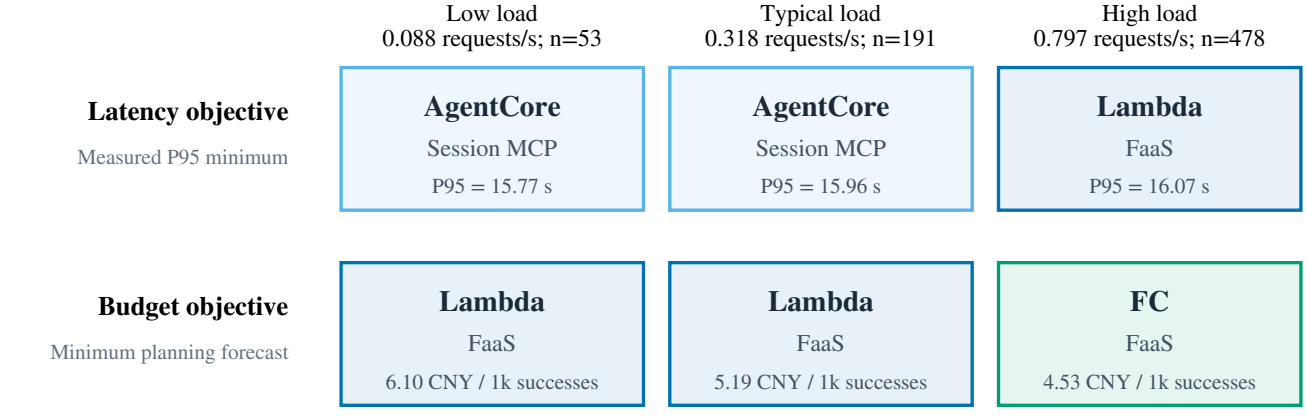

Budget row uses pre-run reservations / observed successes. It is not an actual-cost optimum.

Planning-cost sensitivity

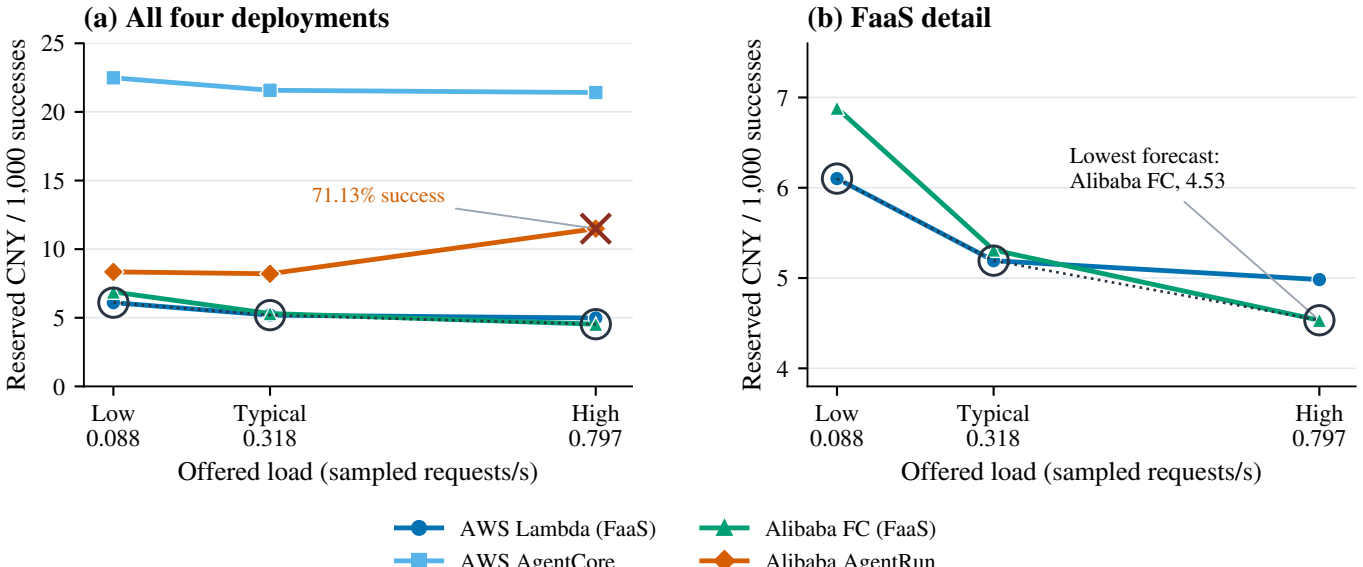

Latency versus planning budget
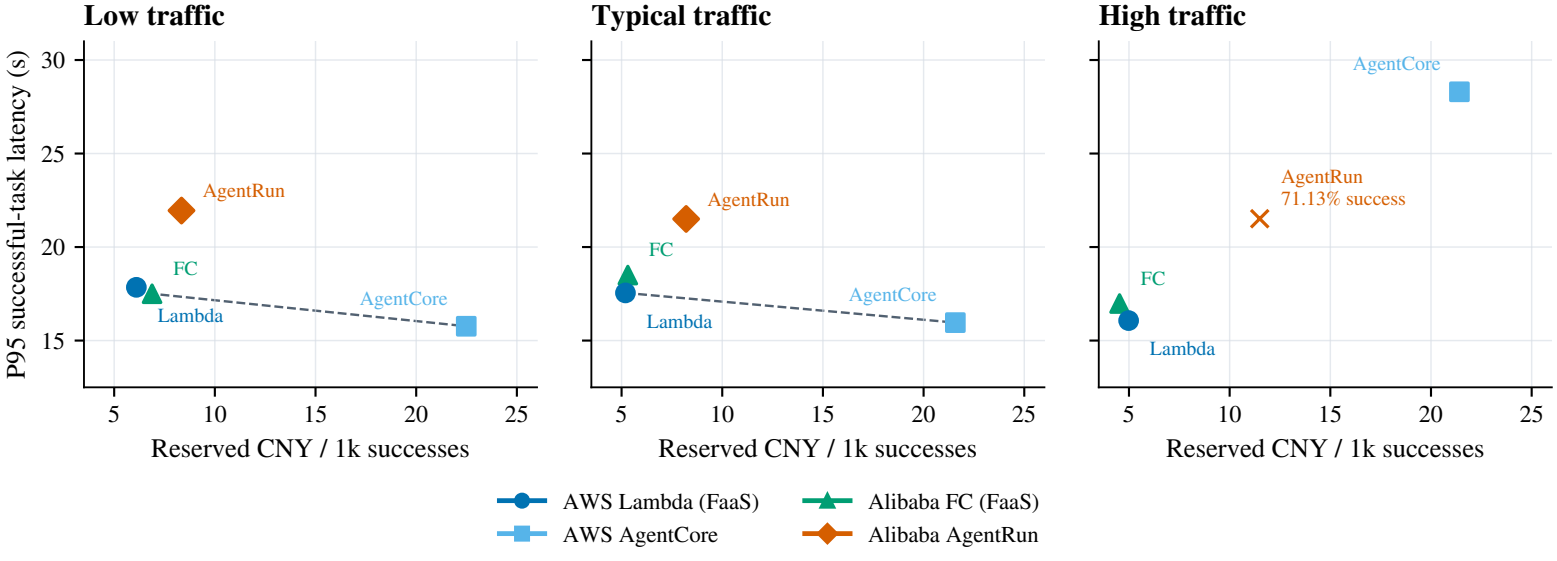
Measured latency; forecast cost. Dashed segments connect feasible nondominated options.